

Tasks and Intertask Communication

Introduction

Multitasking / Multithreading system

Supports multiple tasks

As we've noted

Important job in multitasking system

Exchanging data between tasks

Synchronizing tasks

Sharing resources

Let's now examine these issues

Interprocess / Interthread Communication

When threads operating independently

Our systems have few if any

Conflicts

Chances for corruption

Contentions

Real systems

The interesting ones

Must deal with all such problems

Resources and inter thread communication

Must take place in robust manner

Interaction may be

Direct or indirect

Must be synchronized and co-ordinated

Want to prevent race conditions

Outcome of task or computation

Depends upon order in which tasks execute

Let's begin by looking at shared information

Can occur in a variety of ways

Shared Variables

Simplest solution is shared memory environment

Global Variables

Simplest and fastest of these is

Global variables

Obvious problems

Higher priority process can pre-empt

Modify global data

Shared Buffer

Scheme says two processes share common set of

Memory locations

Producer

Puts data into buffer

Consumer

Removes

Several obvious problems

Arise if one process faster than other

Buffer size critical to avoid such problems

Shared Double Buffer

Scheme says two processes share two common sets of

Memory locations

Called ping-pong buffering scheme

Effective between processes running at different rates

One buffer being filled while other being emptied

Consumer blocks on lack of data

Producer must still avoid over running buffer

Ring Buffer

Scheme FIFO structure studied earlier

Permits simultaneous input and output

Using head and tail pointers

Must be careful to manage

Overflow

Underflow

Mailbox

- Mutually agreed upon memory location

- Two or more tasks use to pass data

- Tasks rely on main scheduler to permit access

- Post* operation for write

- Pend* operation for read

- Pend operation different from poll

- Poll task continually interrogates variable

- Pend task suspended while data not available

- Variety of things passed

- Single bit

- Flag

- Single data word

- Pointer to data buffer

- In most implementations

- Pend operation empties mailbox

- If several tasks pending on flag

- Enabled task resets flag

- Blocks multiple accesses to resource

- On single flag

- Some implementations

- Permit queue of pending elements

- Rather than single entry

- Such scheme may be useful

- Multiple independent copies of critical resource

Messages

- Message exchange is another means for communication

- Now starting to move more into distributed systems

- Called *interprocess communication* facility (IPC)

- Note IPC is not mutually exclusive with shared memory

- Idea to permit processes to communicate

- Without resorting to shared variables

- Particularly in different address spaces

IPC provides two operations

Send

Receive

Messages may be fixed or variable size

Basic Structure

If processes P1 and P2 wish to communicate

Must

Send and receive messages

Establish a communication link

Questions

Variety of questions one may ask

How to establish link

Can link be associated with multiple processes

How many links between pair or process

What is link capacity and are there buffers

What is message size

Are links

Unidirectional

Bi-directional

Implementation methods

Direct / Indirect communication

Symmetric / asymmetric communication

Auto or explicit buffering

Send by copy or reference

Fixed or variable sized messages

Let's look at several of these

Communication

Direct

Each process must explicitly name sender / receiver of message

Messages logically of form

send (P1, message) // send message to P1

receive (P2, message) // receive message from P2

Link properties

Link automatically established between every pair of processes

Processes need only know each others identity

Link associated with only two processes

Between each pair

Only single link

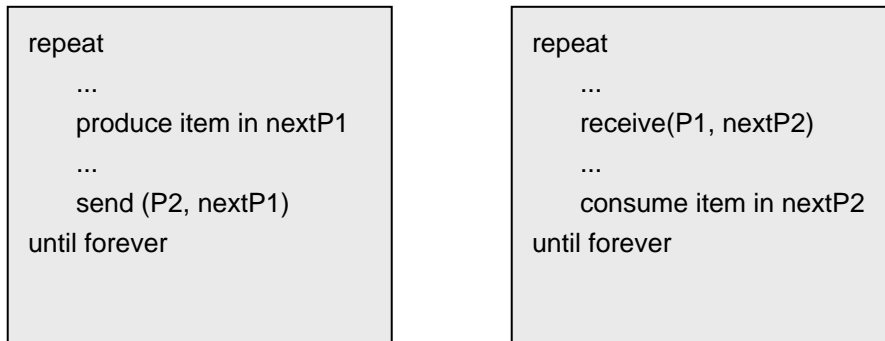
Link may be

Uni/bi directional

Example

Consider skeletal structure

Between producer P1 and consumer P2



Observe scheme uses

Symmetrical addressing

Sender and receiver must name each other

If want asymmetric; addressing

Sender only names recipient

Disadvantage

Ties process name to implementation

Indirect

Messages sent / received from shared variable

Generally in form of mailbox

send (M0, message) // send message to mailbox M0

receive (M0, message) // receive message from mailbox M0

Properties

- Link established

 - Only if processes have shared mailbox

- Link may be associated with multiple processes

- May be multiple links between processes

- Link may be uni/bi directional

Consider 3 processes P0, P1, P2

- All share M0

- Let P0 send and P1 and P2 receive

 - Question - who gets message

Solution

- Associate link with at most 2 processes

- Allow only one process to receive at a time

- Let system select receiver

Mailbox owner

- Process

 - If process owns mailbox

 - Can distinguish between

 - Owner

 - Who can only receive

 - User

 - Who can only send

 - Since each mailbox has unique owner

 - No ambiguity

- System

 - Exists independent of any process

 - OS provides mechanism for process to

 - Create new mailbox

 - Send / receive messages through mailbox

 - Destroy mailbox

 - Creating process

 - May pass access privileges

 - Share mailbox

 - Must manage memory associated with mailboxes

 - For which no process has access rights

Buffering

Establishes number of messages

Temporarily reside in link

Three possibilities

Zero capacity

Link cannot store message

Sender must wait for receiver to accept message

Called rendez vous

Bounded capacity

Message queue has length n

If space remaining

Sender can place message in queue

Continue

Else

Sender must wait for space

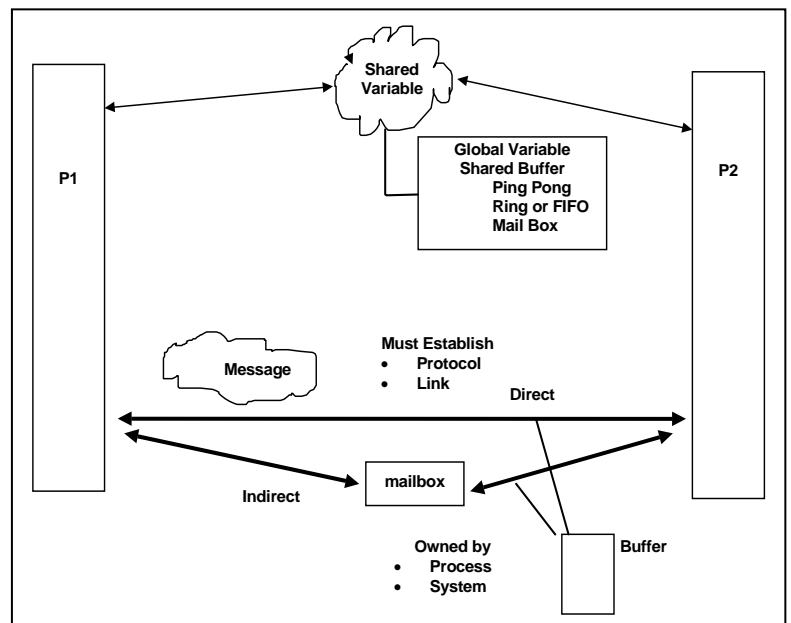
Unbounded

Potentially infinite length

Sender can post message

Continue

No wait



Thread Synchronization

Co-operating threads

One that can affect or be affected by another threads

May directly share logical address space


Code and data

Be allowed to share data only

Through files

Concurrent access to shared data

Can result in data inconsistency



Critical Sections

Consider following problem and code fragments

Exchanging messages through bounded buffer

Allow n items in buffer

Algorithm says

```
Producer
  If not full
    add item
    increment count
  else
    wait until space
```

```
Consumer
  If item
    get item
    decrement count
  else
    wait until item
```

```
Producer
repeat
  ...
  produce an item in nextP1
  while (count == n); // buffer full

  buffer[in] = nextP1;
  in = (in + 1) % n;
  count++;
until forever
```

```
Consumer
repeat
  while (count == 0); // buffer empty

  nextP2 = buffer[out];
  out = (out+1) % n;
  count--;
  ...
  consume nextP2;
  ...
until forever
```

Problem

Value of count

Depends upon who accesses variable

May be any of 3 different values

Variable count is critical variable

Within P1 or P2

Denoted *critical section*

Critical section in general

Section of code in which process is changing common variables

File

Table


etc

While process in critical section

Want to prevent access by all other processes

Termed *mutual exclusion*

Abstractly may represent code as



Semaphores

Solution to critical section problem must satisfy following requirements

mutual exclusion

If process P1 is in critical section

No other process may enter

progress

If no process in critical section and some process wish to enter

Only processes not in remainder section

Can participate in decision

Decision cannot be postponed indefinitely

bounded waiting

Must be bound on number of times other processes can enter critical section

After a process has made a request to enter

Before request granted

Methodology to protect critical section suggested by Dijkstra

Called *semaphore*

Semaphore

Integer or Boolean variable - S

Accessed only through two atomic operations

wait - P(S)

signal - V(S)

Operations may be defined by following code fragments

```
wait(s)
{
    while (s);
    s = TRUE;
}
```

```
signal(s)
{
    s = FALSE;
}
```

s is initialized to FALSE

These may now be used as

```
Process 1
{
    ...
    wait(s)
        critical section
    signal(s)
    ...
}
```

```
Process 2
{
    ...
    wait(s)
        critical section
    signal(s)
    ...
}
```

Consider two concurrently running processes p1 and p2 let

p1

Contain statement s1

p2

Contain statement s2

We require s1 be executed before s2

Thus define semaphore sync

Initialize sync to TRUE

```
p1
...
s1
signal(sync)    // signal
...
```

```
P2
...
wait(sync)      // wait
s2
...
```

Observe

Because synch initialized to TRUE

p2 will execute s2 only after
p1 executes s1

Spin Lock

Main disadvantage of semaphores as described

When wait encountered

Encountering process blocked

Must loop continuously while waiting

Called *busy waiting*

Waiting processes waste CPU cycles while waiting

Other process could use productively

Such a semaphore called *spinlock*

Because process spins while waiting for lock

Advantage of spinlock

No context switch

Can take long time

If lock expected to be held for short time

Spinlock useful

To overcome need for busy waiting

Modify definition of semaphore operations

When process executes wait operation

If semaphore TRUE

Must wait

Rather than wait process can *block* itself

Block operation places self in waiting queue

Associated with semaphore

Process state changed to waiting

Control transferred to scheduler

Blocked process should be restarted

Some other process executes signal operation

Process

Restarted

By *wakeup* operation

Places process in ready state

Placed in ready queue

Semaphore now defined as follows

s initialized to 1

```
wait(s)
{
    s = s-1; // on first pass s == 0
    if (s < 0)
    {
        add process to waiting queue;
        block;
    }
}
```

```
signal(s)
{
    s = s+1;
    if (s <= 0)
    {
        remove process from waiting queue;
        wakeup(p);
    }
}
```

Note semaphore now has integer value

block operation suspends invoking process

wakeup resumes execution of blocked process

Both operations provided by operating system calls

Observe

Waiting list can be implemented by linked list

Perhaps implement as FIFO queue

Mutexes and Counting Semaphores

Semaphores we've looked at called *binary semaphores*

Can take on either one of two values

Mutex

Binary

Used to serialize access to reentrant code

Allows only one thread into controlled code section

Example

Key to toilet

Semaphore

Counting

Can take on more than two values

Like previous example

Used to protect pools of resources or track number of resources

Restricts number of simultaneous users (threads) of shared resource

Example

Number of keys to toilet

Working with a counting semaphore - let's call these

wait - wait(s)

signal - sig(s)

```
wait(s)
{
    s--;
    if (s<0)
        add this process to queue;
    block;
}
```

```
sig(s)
{
    s++;
    if (s<=0)
        remove a process Pi from queue;
        wakeUp(Pi);
}
```

Each semaphore has

Integer value

List of associated processes

When process must wait on semaphore

Added to list of processes

Signal

Removes process from list

Awakens it

Operations may be defined by following code fragments

Bounded Buffer Problem

Let's look at one classic synchronization problem

Consider we have a pool of n buffers

Each can hold one item in this example

We define semaphores

mutex

Provides mutual exclusion for accesses to buffer pool

Initialized to value 1

Empty - semaphore

Count number of empty buffers

Initialized to n

Full - semaphore

Count number of full buffers

Initialized to 0

Code fragments illustrated as

Producer

```
repeat
  ...
  produce an item anItem
  ...
  wait(empty); // check for non zero
                // dec empty cnt
  wait(mutex);
  ...
  add anItem to buffer nextProd;
  ...
  signal(mutex);
  signal(full); // inc full cnt
  ...
until false
```

Consumer

```
repeat
  wait(full); // check for non zero
                // dec full cnt
  wait(mutex);
  ...
  remove anItem from buffer nextCons;
  ...
  signal(mutex);
  signal(empty); // inc empty cnt
  ...
  consume item anItem
  ...
until false
```

Readers and Writers Problem

Data object may be shared among several concurrent processes

Some may want to read and others may want to write

Processes referred to as

Readers

Writers

If 2 readers access simultaneously

No problem

If writer and any other process access simultaneously

Big problem

Referred to as *readers - writers* problem

Several variations

First readers-writers

No reader waits

Unless writer has obtained access of shared variable

Second readers-writers

Once writer ready

Performs write as soon as possible

If writer waiting

No new reader started

Solution to first readers-writers problem

Define

Semaphores - mutex, wrtSem

Initialize to 1

mutex - ensure mutual exclusion when readcount updated

wrtSem - mutual exclusion for writers

integer - numReaders

Initialize to 0

numReaders - count of readers currently accessing shared variable

Code fragment given as:

Writer Process

```
wait(wrtSem);           // wait for wrtSem == 1
                        // wrtSem = 0

...

perform writing;

...

signal(wrtSem);         // wrtSem = 1

...
```

```

Reader Process
wait(mutex);           // wait while mutex == 1
                        // mutex = 0

numReaders++;          // inc number of readers
if (numReaders == 1)   // if i'm the only reader
    wait(wrtSem);      // make sure no writers
                        // wrtSem = 1

signal(mutex);         // mutex = 0

...
Perform reading;
...
wait(mutex);           // wait for mutex == 1
                        // mutex = 0

numReaders--;          // dec number of readers
if (numReaders == 0)   // no readers
    signal(wrtSem);    // wrtSem = 0
signal(mutex);         // mutex = 0

...

```

Note

If writer in critical section and n readers waiting

One reader queued on wrtSem

n-1 readers queued on mutex

If writer executes signal(wrtSem)

May resume

Waiting readers

One waiting writer

Decision made by scheduler

Monitors

Semaphores we've studied

Fundamental synchronism mechanism

However low-level mechanism

Easy to make errors with them

Monitors are program modules

Offer more structure than semaphores

Implementation can be as efficient

Monitors

Data abstraction mechanism

Encapsulate

Representation of abstract object

Provide public interface

Only means by which

Internal data may be manipulated

Contains variable to

Store object's state

Procedures that implement operations on object

We satisfy mutual exclusion

By ensuring

Procedures in same monitor

Cannot execute simultaneously

Conditional synchronization

Provided through condition variables

Monitor used to group

Representation and implementation

The interface and body

Of shared resource

Has *interface* and *body*

Interface

Specifies operations and behaviour provided by resource

Body

Contains

Variables

Represent state of resource

Procedures

Procedures

Implement operations specified in interface

Schematically we have

```
monitor monName
{
    initialization statements //analogous to constructor
    procedures
    permanent variables
}
```

Procedures implement

- Visible operations

Permanent variables

- Shared by all processes

- In the monitor

- Like statics in C++ or pool variables in Smalltalk

- Denoted permanent

- Retain values on exit

- As long as monitor exists

Procedures

- May have local variables

By virtue of being an Abstract Data Type

Monitor is a distinct scope

- Only procedure names – this is the public interface

- Visible outside of monitor

Permanent variables

- Can only be changed

- Through one of the visible procedures

Statements within monitor

- Cannot affect variables outside monitor

- In different scope

Permanent variables

- Initialized before any procedure called

- Accomplished by

- Executing initialization procedures

- When monitor instance created

Monitor sounds very similar to C++ class

Major difference

Monitor shared by multiple concurrently executing processes or threads

Consequently

Threads or processes using monitor

May require

Mutual exclusion

To monitor variables

Synchronization

To ensure monitor state conducive to continued execution

Mutual exclusion

Usually implicit

Synchronization

Implemented explicitly

Different processes require different forms of synchronization

Implementation achieved through

Condition variables

Shared variables discussed earlier

Monitor procedure

Called by external process or thread

A procedure is active

If a thread or process executing

Statement in procedure

At most one instance of monitor procedure

Active at any one time

Cannot have

Two different procedures invoked

or

Two invocations of same procedure

By definition

Execute with mutual exclusion

Ensured by

Language

Library

Operating system

- Generally implemented
 - Locks or semaphores
 - Inhibiting certain interrupts

Condition Variables

- Condition variables used as part of synchronization process
 - Used to delay thread or process that
 - Cannot safely continue
 - Until monitor's state satisfies some Boolean condition
 - Used to awaken delayed process
 - Once condition becomes true

Condition variable

- Instance of variable of type *cond*
 - `cond myCondVar;`
- Can only be declared inside monitor
- Value of condition thus it represents a queue
 - Queue of delayed processes
- Initially queue is empty
- Value can only be accessed indirectly
 - Much like private variables in C++ or Java
- Test state
 - `empty(myCondVar);`
- Thread can block on a condition variable
 - `wait(myCondVar);`
 - Execution of wait causes process to
 - Move to rear of queue
 - Relinquish exclusive access to monitor
- Blocked process awakened
 - `signal(myCondVar);`
 - Execution of signal causes thread
 - At head of queue to awaken

Execution of signal

- Seems to cause dilemma
 - Upon execution two processes have potential to execute
 - Awakened thread

Signaling thread

Contradicts requirement

Only single thread active in monitor at once

Two possible paths for resolution

- Signal and continue

Signaling thread continues

Awakened process resumes at some delayed time

Considered nonpreemptive

Process executing signal

Retains exclusive control of the monitor

- Signal and wait

Considered to be preemptive

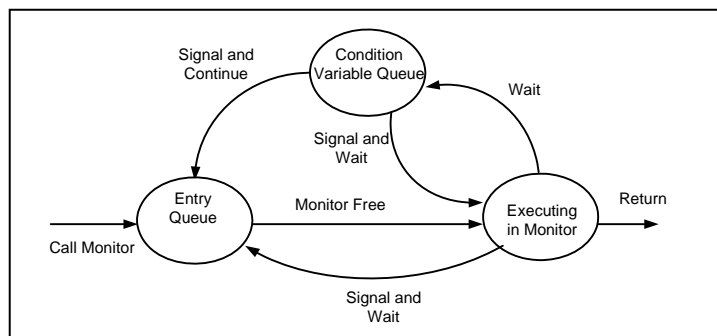
Process executing signal

Relinquishes control and passes lock

To awakened process

Awakened process preempts signaling process

Can describe process with following state diagram



Operation / synchronization occurs as follows

Thread *calls* monitor procedure

If another thread executing in monitor

Caller placed into *entry queue*

When monitor becomes free

Result of *return* or *wait*

One thread moves from entry queue into monitor

Else passes through entry queue

Begins executing immediately

If thread executes wait on a condition variable
While executing in monitor
Thread enters queue associated with that variable

When thread executes
Signal and Continue on a condition variable
Thread at head of associated queue
Moves to entry queue
Signal and Wait on a condition variable
Thread at head of associated queue
Moves to monitor
Thread executing in monitor
Moves to entry queue

Bounded Buffer Problem with Monitor

Let's look at one classic synchronization problem
Looked at earlier with semaphores
Implemented with monitor
Consider we have a pool of n buffers
Each can hold one item in this example

We define a monitor *boundedBuffer*

We define condition variables

notEmpty

Signaled when buffer count > 0
Tracks empty buffers
initialized to 0

notFull

Signaled when buffer count < n
Tracks full buffers
Initialized to 0

We define procedures

put(data)

Puts data into a buffer

When space available

get(data)

Gets data from a buffer

When data available

We define the protected entity

bufferPool

We can implement our monitor as follows

```
monitor boundBuffer
  bufferPool;
  count = 0;
  cond notEmpty;      // signaled when count > 0
  cond notFull;       // signaled when count < n

  put(anItem)
  {
    while(count == n) wait (notFull);
    put anItem into a buffer
    signal (notEmpty);
  }
  get(anItem)
  {
    while(count == 0) wait (notEmpty);
    get anItem from a buffer
    signal (notFull);
  }
}
```

Code fragments illustrated as

```
Producer
repeat
  ...
  produce an item anItem
  ...
  boundBuffer.put(anItem)
  ...
forever
```

```
Consumer
repeat
  ...
  boundBuffer.get(anItem)
  ...
  consume item anItem
  ...
forever
```

Deadlocks and Starvation

Deadlocks

Implementation of semaphore or monitor with waiting queue

Can result in situation in which 2 or more processes

Wait indefinitely

Called *deadlock*

Consider 2 processes P0 and P1

Let each process have 2 semaphores

S1 and S2

May be resources each needs

R1 and R2

Need both to continue

Let R1 and R2 be set to value 1

Let

P0 set wait(S1) // wait for R1 decrement S1 (=0)

P1 set wait(S2) // wait for R2 decrement S2 (=0)

Now let

P0 set wait(S2) // wait for R2 decrement S2 (= -1)

P1 set wait(S1) // wait for R1 decrement S1 (= -1)

At this point

P0 must wait for signal(S2)

P1 must wait for signal(S1)

These operations cannot be executed

Processes blocked

Every process in set waiting for event

Possible only by another member in set

Will discuss in much greater detail shortly

Starvation

Problem called *starvation* can occur

Process waiting within semaphore

Other processes added or removed

LIFO order

Events and Signals

- Some languages provide mechanisms for handling

 - Asynchronous events

- Provides software interrupt

 - Generally used for exceptions

 - Divide by zero

 - Arithmetic overflow

 - etc.

- In addition to built in procedures

 - Some permit user defined procedures to be

 - Provided and executed

 - ANSI-C

- Provides signal and raise

 - Signal

 - Software interrupt handler

 - Responds to exceptions indicated by raise

 - Raise

 - Mechanism to signal an exception or event

- Both implemented as function calls

- Passing pointers to functions

 - Can handle variety of events or exceptions