

Deadlocks

Introduction

In multitasking or multiprogramming system

Several processes can compete for finite number of resources

A process requests resources

If not available process placed in wait state

If resources never become available

Process remains waiting

Called *deadlock*

Illustration

Law passed in Kansas around turn of century

When two trains approach each other at crossing

Each shall come to full stop

Neither shall start up again until the other has gone

We talked about deadlocks earlier

Will now look at methods to deal with problem

Note

Most contemporary operating systems

Do not use deadlock prevention techniques

As systems become more complex and number of processes increases

Problem will need to be addressed

System Model

System has finite number of resources

These distributed among number of competing processes

Resources partitioned into several types

Identical resources

May have multiple instances of same resource

Printers

Telecomm channels

Memory space

Allocation of any one may be sufficient

Note printers may be identical

If convenience to user compromised

May not be considered identical

Printer on 1 and 9th floors of office building

Dissimilar resources

Second kind of resource

Those that are unique for one reason or another

Single copy

Identical printers for example may not be identical

If convenience to user compromised

May not be considered identical

Printer on 1 and 9th floors of office building

To use resource process must request resource
May request as many resources as it wishes
May not exceed total number

Under normal operation may only utilize resources in following order

Request

If can't be granted immediately
Requesting process must wait

Use

Process operates on or uses resource

Release

When finished give up resource

Request and release are system calls

Deadlock Characterization

Set of processes in deadlock state
Every process in set is waiting for event
Can be caused only by another process in set
Events of concern
Resource acquisition and release

Necessary Conditions

For deadlock to occur following must hold simultaneously
Note these are necessary not sufficient
Note also these are not independent

Mutual Exclusion

A least one resource held in non-sharable mode

Hold and Wait

Must be process holding resource and waiting for additional
Being held by other processes

No Preemption

Resources cannot be preempted

Circular Wait

Set of processes $\{P_0 \dots P_n\}$ such that
P0 waiting for resource held by P1
P1 waiting for resource held by P2

...

Resource Allocation Graph

Can understand deadlock formally using
Resource Allocation Graph
Resource allocation graph
Directed graph
Set of

Vertices V

Partitioned into two sets

Set of processes $\{P_n\}$

Set of resources $\{R_m\}$

Edges E

Connecting

$\{P_n\}$ to $\{R_m\}$

$\{R_m\}$ to $\{P_n\}$

Directed edge from P_i to R_j

$P_i \rightarrow R_j$

Signifies P_i requested R_j and is currently waiting

Called *request* edge

Directed edge from R_j to P_i

$R_j \rightarrow P_i$

Signifies R_j allocated to P_i

Called *assignment* edge

Graphically

Process

Circle

Resource

Rectangle

Multiple copies

Signified by dot in rectangle

Consider following RAG


We have the following situation

Sets

$P = \{P_1, P_2, P_3\}$

$R = \{R_1, R_2, R_3\}$

$E = \{P_1 \rightarrow R_1, P_2 \rightarrow R_2, R_1 \rightarrow P_2, R_2 \rightarrow P_3, R_3 \rightarrow P_1, R_3 \rightarrow P_2\}$



Resource Instances

1 of R1

1 of R1

2 of R3

3 of R4

Process States

P1

Holding 1 R3

Waiting for R1

P2

Holding 1 R1 and 1 R3

Waiting for R2

P3

Holding 1 R2

Using techniques from graph theory

Can show if contains no cycles

No process in system is deadlocked

If cycle exists

Potential for deadlock exists

Does not guarantee

If single instance of each resource

Cycle

Implies deadlock has occurred

Becomes necessary and sufficient condition

If multiple instances

Cycle


Does not necessarily imply deadlock

Necessary but not sufficient condition

Let's look at two examples

First has cycle and deadlock

Second has cycle and no deadlock



Handling Deadlocks

Let's now look at some ways of dealing with

Deadlock problem

Several ways

Use protocol to ensure deadlock will never happen

Allow system to enter deadlock state and recover

Ignore problem

Solution used by most operating systems

Including UNIX

Ensuring No Deadlock

Can use

Deadlock prevention


Ensure one of necessary conditions cannot occur

Deadlock avoidance

Requires additional information

Which resources process will request and use

During lifetime



Deadlock Prevention

Let's look at prevention first

Easiest solution

Mutual Exclusion

Must hold for non sharable resources

Single printer

Sharable resources

Mutual exclusion not required

Read only files

Cannot prevent deadlocks by denying mutual exclusion

Some resources inherently non-sharable

Hold and Wait

To prevent hold and wait condition

Must guarantee when process requests resource

Does not hold any other resources

Protocol 1

Request and be allocated all resources

Before execution

Protocol 2

Can only request resources when have none

Can request resources and be allocated

To request additional

Must give up what have

Two main disadvantages

Resource utilization low

Allocated but not used for long time

Starvation possible

Process needing popular resources may have to wait indefinitely

No Preemption

To prevent no preemption condition

Protocol 1

If holding resources and need more that are not available

Process must wait

All resources currently being held

Preempted

Added to list of resources for which process is waiting

Process restarted when it can

Regain old resources

Acquire new ones it requested

Protocol 2

If process requisite resources

Check
 If available
 Allocate
 else if with another process waiting for resources
 If with another waiting process
 Preempt
 Allocate to requesting process
 else
 wait

Circular Wait

To prevent circular wait
 Place total ordering on all resources
 Require each process to request resources
 Increasing order of enumeration
Let $R = \{R_1, R_2, \dots, R_m\}$ be set of resource types
 Assign each type unique integer number
 Allows ordering relation to be applied and evaluated
Protocol 1
 Initially request any desired resources
 Additional resource requests
 Only in increasing order of enumeration
 If multiple copies of single resource needed
 Must request all at once
Protocol 2
 Initially request any desired resources
 Additional resource requests
 If request R_j
 Must release any resources $\{R_i\}$ such that $i \leq j$

Deadlock Avoidance

Deadlock prevention algorithms
 Prevent deadlocks by restraining requests
 Restrains ensure
 At least one of necessary conditions cannot occur
Consequence
 Low utilization of resources
Deadlock avoidance
 Requires additional information
 About how resources requested
 Consider system
 Resources
 Printer
 Tape drive
 Processes

P1 and P2
Need
P1
Printer then tape drive
P2
Tape drive then printer
Knowledge
Knowing need in advance
Permits scheduling to ensure no deadlock

Various algorithms
Require differing amounts of information
Let's walk through simple one to get idea

Declare in Advance

Simplest most useful model requires each process to declare
In advance
Maximum number of resources
Of each type it may need
Given such information
Possible to construct algorithm
To ensure system will never enter deadlock state
Such a scheme defines basis for deadlock avoidance
Avoidance algorithm
Examine resource allocation state
Defined by number of
Available and allocated resources
Max number of demands by processes

Safe State

Resource allocation state is *safe*
System can
Allocate resources to each process
In some order
Avoid a deadlock


Formally

System is in *safe state*
If there exists a *safe sequence*
Sequence of processes $\langle P_1, P_2 \dots P_n \rangle$ is safe sequence
For current allocation state
If
For each P_i
Resources that P_i can still request can be satisfied by
Currently available resources plus
Resources held by all P_j such that $j < i$
Observe if needed resources not available

P_i can wait until P_j have finished
 Can then have all needed resources
 When P_i finishes P_{i+1} obtain needed resources
 If no such sequence exists
 System state is unsafe

Observe

Safe state is not deadlock state
 Deadlock state is unsafe state
 Not all unsafe states are deadlock states
 Unsafe state may lead to deadlock
 Three spaces illustrated as



Example

Consider following system

12 I/O Ports

3 Processes

Let max and current needs be given as

	Max Needs	Current Needs
P0	10	5
P1	4	2
P2	9	2

We have total allocation of 9 with 3 ports free

At time t₀

System in safe state

Sequence <P₁, P₀, P₂>

Safe sequence

Can satisfy P₁

Block P₀

Until P₁ finished

Block P₂

Until P₁ finished

At time t₁

System can go to unsafe state

Let P₂ requests additional port

Only P₁ can be allocated all resources

When it returns them

Only 4 total available

P₀ allocated 5 ports

Max need of 10

May request 5 more

Not available so block
P2 may request additional 6
Not available so block
Deadlock

Avoidance Algorithms

Resource Allocation Graph Algorithms

If we have system with one instance of each resource

Can use variant on resource-allocation graph to avoid deadlocks

Introduce new edge type - *claim* edge

Claim edge $P_i \rightarrow R_j$

Indicates Process P_i may claim resource R_j sometime in future

Edge has semantics similar to request edge

Direction same

Notation is dashed line

Requires that resources be claimed a priori in system

Before process starts executing

All claim edges must be present in resource-allocation graph

Restriction may be relaxed to allow addition of claim edge

If all other edges from process are claim edges

Protocol

When P_i requests R_j

Claim edge converted to request edge

Similarly when resource R_j released by P_i

Request edge converted to claim edge

Claim edge can only be converted into request edge

If conversion does not result in cycle

If no cycle exists

Allocation will leave system in safe state

Observe

If P_2 requests and is allocated R_3


Although available

Cannot allocate

Will create cycle and thus unsafe state

If P_1 requests R_3

We have a deadlock



Deadlock Detection

If system does not employ prevention or avoidance algorithm

Deadlock may occur

In such environment

System must provide

Algorithm to determine if deadlock has occurred

Algorithm to recover from deadlock

Detection in Single Instance Environment

As with avoidance

Can use variation on resource allocation graph

Called *wait-for* graph

Algorithm



Start with resource-allocation graph

Remove nodes of type resource

Collapse appropriate edges

Will result in graph with only processed

Deadlock exists if and only if the graph contains a cycle



Deadlock Recovery

When deadlock algorithm detects deadlock exists

Several possible alternatives

Inform user

Difficult in embedded system

Let system recover automatically

Automatic recovery

Two general schemes

Abort

All processes

One at a time

Preempt resources

Process Termination

All deadlocked processes

Will clearly break deadlock

At great expense

Processes may have computed for long time

All results may be lost

One process at a time

Until deadlock cycle eliminated

Involves considerable temporal overhead

As each process aborted

Must rerun deadlock detection algorithm

Extreme care must be taken

Aborting process may leave resources in

Unknown or unusable state

Must also determine which process to abort

Similar to CPU scheduling problem

Want to abort processes in terms of increasing cost

Potential factors

Process priority

Time since start and remaining run time

Resource mix and quantity

Resource demand to complete

Number of processes to be terminated

Resource Preemption

Method requires

Successive preemption of resources

Allocation to other processes

Until deadlock cycle broken

If preemption used

Must consider three issues

1. Selecting a victim

Must determine order of preemption to minimize cost

Factors include

Number of resources deadlocked process holding

Amount of elapsed execution time for deadlocked process

2. Rollback

If resource preempted

What should be done with associated process

Cannot continue

Often cannot determine completely safe state

Simplest solution is complete rollback

Abort process and restart

Can try to roll back as far as necessary to break deadlock

Entails maintaining information on all running processes

3. Starvation

How to ensure starvation will not occur

Want to ensure resources not always preempted from same process

Summary

Deadlock occurs when two or more processes

Waiting for event that can only be caused by one of waiting processes

3 major methods for addressing

Use protocol to ensure will never enter deadlock state

Allow system to enter deadlock state and recover

Ignore problem

Deadlock can only occur

If and only if 4 conditions occur simultaneously

We prevent deadlock

Ensuring one condition will not occur

If system does not employ protocol to ensure deadlock does not occur

Then detection and recovery scheme must be employed

If deadlock detected

Can recover by global or selective termination

Process

Resources