

Interfacing to the Outside World

Introduction

There are variety of ways and reasons to talk to outside world

Consider first kinds of devices we may want to interface to

For embedded systems tasks include

- Measuring
- Controlling
- Interacting with other subsystems
- Computing
 - Accelerators
 - Graphics
- Communications

Collectively such devices called I/O system

Devices we may interact with include

File systems and storage devices

Read and write data

Network

Send and receive data

Execute remote procedures

Keyboard and mouse

Interact with computer

Scanning devices

Bring in information

Display and printing devices

Display variety of data

We interact with external devices in several ways

Busses

Comprised of

Address

Control

Data

Serial

Bit and character

Character

Parallel

Character

Character and Word

Timing

Asynchronous

Transfer independent of sender / receiver timing

Synchronous

Transfer dependent upon sender / receiver timing

Busses a First Look

To understand input and output

We must have some understanding of

How we communicate with such devices

Three general configurations

Star

Ring

Bus

Each has certain advantages and disadvantages

For all configurations

Two kinds of messages

Addressed to single device

Multicast

Star

Master slave kind of arrangement

Device to device communication must go through master

Transmit

Device at center directs activities and message exchange

With all other devices

Receive

Master transmits to desired destinations

Failure

If one device or link fails

Communication with others continues

Ring



Variants are common in communication networks

Basis for token ring networks

Embedded processors in an automobile

There is typically no bus master

Device accepts all messages circulating in ring



If device receiving message ADDRESSED device

Message accepted

Else

Message passed on o next device

Transmit

Some protocol used to decide who is able to transmit messages

To use the bus

Receive

All devices can listen for message

Failure

Ring usually implemented as two concentric rings

If device fails or ring severed

Can repair itself by reconfiguring

Bus

Variation on the star architecture

May or may not be a bus master


Simple bus probably one of more common architectures

Device interconnection several ways

All can transmit or receive

Some transmit or receive only

Control bus to all or request / grant configuration



Transmit

If bus master

It addresses device

Sends information

Permits another device to transmit

Else

Some protocol used to decide who is able to transmit messages

To use the bus

Messages may be

- Broadcast
- Multicast
- Polled

Receive

All devices can listen for message and act on message

Failure

Individual device failure will not compromise net

Severed net can prevent communication beyond severed point

Whichever scheme we choose

Generally must provide three basic functions

- Data
- Address
- Control

Address

An important question

How do we define and establish an address

Initial address for a device

Subsequent addresses as device added to system

Current schemes

Geographic addressing

Addresses are dynamically assigned during initial configuration


As new device attached to system and detected

Forward capabilities

Serial System

Address sent down same path as data

Appears first to select receiver



Typically address and data

Have some means to distinguish

Bit pattern

Parallel System

Address may be handled in two main ways


Transmit over same bus as data

Precedes data

Works like serial method only faster

Tagged to distinguish from data

Either as part of transmission or with control lines



Transmit over separate bus

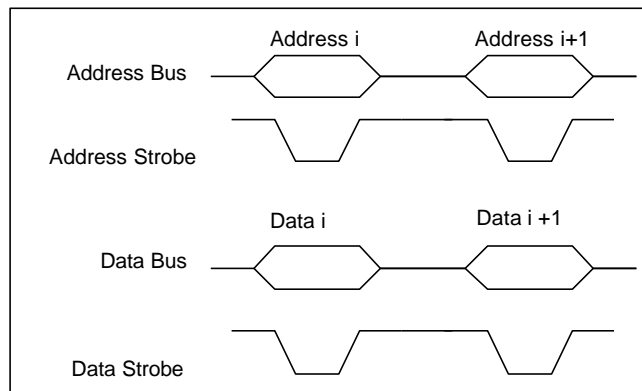
Bus may be same size as data

Depending upon architecture

May be larger if also used to access memory

Smaller is special bus

Transmitted simultaneously with data



Control

Number and structure depend upon nature of system

Typical signals may include

Read / Write

Address / Data present or stable

Clock

Where does the clock come from

Separate line

Encoded in the data

Transmission Direction

Ready or Active

Synchronize

Reset

Power

May or may not be counted as control signal

Data

Whether serial or parallel

Data lines carry information we wish to transmit or receive

When in parallel

Will typically carry one word

For system with 32 bit words

Bus would be 32 bits wide

Carry DB0 - DB31

Important to determine which is the MSB and which the LSB

Moving Outside of the Processor

A First Look

The procedure serves as basis for look at

Interacting with outside world signals

As we've noted embedded systems

Designed to perform some task

Performing task means interfacing with world outside processor

Such tasks may include

Making measurements

Controlling motor

Running telecommunication system

We are dealing with portions of von Neumann machine

Designated

Input and Output

Memory

When consider memory mapped I/O

Our exchange with outside world

Synchronous

Based upon some timing element

Asynchronous

Start of and subsequent transactions temporally independent

When studying processor input and output

We must look at several things

- Data source or destination
 - I/O ports
 - Memory address space
- Nature of the exchange
 - Three types possible
 - Event
 - Shared variable
 - Message
 - Note
 - Such exchange also occurs between / among Tasks within same processor
- How I/O procedure invoked
 - Associated restrictions
- Where procedure resides
- Protocol for the data exchange
- Timing requirements

Let's look at each of these


Data Source or Destination

Data may be exchanged with external world

Two ways

I/O ports

Memory address space



I/O Ports

Input / Output Ports

Number of lines on processor dedicated

Bringing data into processor

Sending data out of processor

Direction

Lines may be

Input

Output

Input and Output

Unidirectional lines tend to be static

Direction established

Hardware

During initialization
Bi-directional lines are dynamic
Direction selected
Based upon nature of transaction
Usually grouped in sets of 4 or 8
Direction for entire group set

Memory Address Space

Read from or written to
Locations within the processors
Primary memory address space
Will elaborate shortly
When discuss memory mapped I/O

Architecture



Let's begin at a high level and examine several I/O structures
Within computer we use two basic approaches
Memory Mapped I/O
Peripheral Processor

Memory Mapped I/O

Basic idea is to for memory and I/O to share the same address space
That is map I/O space into an memory address
I/O reads and writes are done to memory addresses
Traditional scheme used in PC
Such a scheme places the responsibility of external communication
Onto CPU
Simple to implement
Adds extra task

Implementation

Assign address to each input or output device
Some devices may have multiple addresses
Read
Write
Such addresses are typically hard coded
Using switches or jumpers
Early PCs used such a scheme
Expansion led to
Address conflicts
IRQ conflicts



Limited expansion

Advantage

Easy to implement

Low cost

Disadvantage

Extra burden on CPU

Potentially slow

Peripheral Processor

As noted with memory mapped I/O

CPU must be involved in all transactions

Including detailed timing

Can place significant burden on CPU

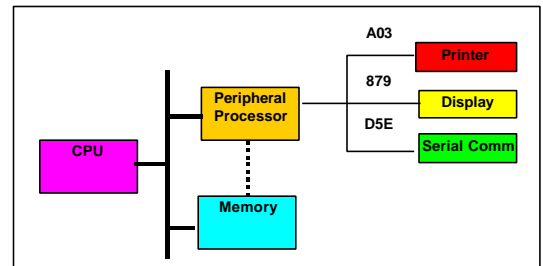
Peripheral processor scheme

Dedicates processor to handle all I/O tasks

Basic architecture appears as

Peripheral processor

May or may not be connected to memory



Several interface schemes utilized

Various levels of involvement of

System Bus

CPU

Non peripheral devices

Alternative protocols

Based upon level of CPU involvement

- CPU send / receive from peripheral processor

PP

Manages all I/O

Uses bus in bursts during exchange

Signals CPU when data available

- CPU tells peripheral processor where to find / put data in memory

PP

Exchanges data with memory

Manages all I/O

Signals CPU when data available

Uses bus

In bursts during exchange

For duration of exchange

DMA

Using address passing and shared buffer

- CPU enables non-peripheral device and peripheral processor to communicate

PP

Exchanges data with device

Manages all I/O

Signals Device when data available

May signal CPU when complete

Uses bus

In bursts during exchange

For duration of exchange

Advantage

I/O speed independent of CPU

Unburdens CPU

Disadvantage

Higher cost

Complex

Nature of the Exchange

Three types possible

Event

Shared Variable

Message

Observe these are abstractions

Within each type

Variety of related signals

Exchange may be used

Outside world

Between tasks within processor

Exchange types

Event

Event is change in state signal

Usually assumed to be single signal

Generally asynchronous to executing process

Every occurrence of event

Simultaneously activates

Functions procedures to tasks

Linked to it

Occurrence of event

May or may not be stored in some way

Becomes important issue in real time systems

Acquired in several ways

Sampled

One example called *polling*

Arrives in form of asynchronous input

One example called *interrupt*

Processor may or may not respond

Generated

Under software control of

I/ O port or line

Write to data line accompanied by control signal

Shared Variable

Variable may be read or written by multiple

Processors

I/O devices

Used to exchange data between asynchronous functions

Because no inter process timing constraints

Integrity of data must be respected

Such shared data

Represent critical section

Protect with

Semaphore

Monitor

As we learned earlier with shared data

Among tasks and threads

Shared variable(s) may be

- Global

Normally discouraged in traditional programming

We use such schemes in real time embedded systems

Eliminate cost of passing parameters via stack
Here we may designate
A global buffer area into which to store data
Pointer to buffer area

- Passed
 - Usually done as pointer to buffer area
 - Passed by task or thread
 - Wishing to do I/O

Direction

Unidirectional

Bi-directional

Modification

Read

Read/Write

Nature may be

Complete word

If in memory space

1 to several bits if in I/O port space

Complex data type

Pointer exchanged

Message

May be exchanged via

Proprietary network

One of several standards

Which ever approach used

Typically implemented as hierarchy virtual networks

Above the hardware

Varying number of software layers or levels

At each level

A different language is spoken

Referred to as a *protocol*

The function is to provide services for the level above

Relationship


Service provider

Service consumer

Entire collection called

A *network architecture*

Set of protocols used by a machine called



A protocol stack

Information sent on each level called

A message or messages

It is possible that message on higher level

Composed of several lower level messages

Two major protocol schemes or stacks used

OSI

Open Systems Interconnection model

Proposed and developed by

The International Standards Organization

Comprises 7 layer virtual machine

TCP /IP

Transmission Control Protocol / Internet Protocol

Comprises 5 layer virtual machine

Physical and data link layers of OSI

Combined into

Host to network layer

OSI	TCP/IP
Physical	Host to Network
Data Link	
Network	Internet
Transport	Transport
Session	Not Present
Presentation	Not Present
Application	Application

When we elect to communicate via either standard

We generally integrate a

Commercially available protocol stack

Messages comprise

Several complete words

Simple

1 - 2 words transferred

Complex

Several blocks of data

Message exchange

Can be viewed as producer - consumer relationship

Usually follows some protocol

Start of message

Header information

Body

Error management

End of message

Message may be transmitted

Serial / Parallel

Serial

Data transferred into out of processor

1 - 2 lines

1 Line

May be

Unidirectional

Input or output only

Half duplex

Bi-directional

One direction at a time

Controlled by protocol

Duplex

2 Lines

Data transferred in both directions

Simultaneously

Full duplex

Perhaps some control lines

Parallel

Usually width of word

Transmission

Unidirectional

Input or output only

Bi-directional

One direction at a time

Timing and Data Exchange

Interface can be designed to be synchronous or asynchronous

Let's look first at the asynchronous


Asynchronous

No clock

Exchanges co-ordinated using some form of handshaking protocol

Such protocols can be simple or complex

Typical examples



Strobe associated with each data word

No acknowledgment of acceptance


Observe

Strobe can be of either polarity

Strobe associated with each data word

Each data word acknowledged

With return strobe



Full handshake


Ready for data

Here's data

I've got it

OK

Exchange looks like this



Time from bit transition


Often used in serial protocols

Agreed upon

Change in state on data line signals

Start of transmission

End of transmission



Potential problems

Difficult to test

Clock noise more difficult to filter out

Potentially complex protocol

To identify start / end of transmission

Potential advantages

Devices may run at different / differing speed

No clock skew on long busses

Synchronous

There are several drawbacks of the asynchronous transmission schemes

Extra overhead of control bits

Bit clock synchronization scheme less reliable at higher data rates

Problems can be alleviated with synchronous transmission

Still must achieve

Bit, character, frame synchronization

Frame synchronization usually derived from the former
Generally includes clock in control lines

Exchanges between sender and receiver

Synchronized to the clock

- Directly
- Signals derived from the clock
Manchester phase encoding

Serial Exchange

Clock either separate or encoded in data

Parallel Exchange

Clock one of control lines

Bit Synchronization

Two schemes typically used

Encode the clock in the data

Re-derive the clock from the data

Encoded Clock

Three different methods generally used

Bipolar Encoding

Binary 0's and 1's are represented by

Different polarity signals


Each bit cell contains clocking information

Observe

Signal returns to zero level after each encoded bit

Referred to as *return-to-zero RZ* signal

Scheme requires 3 distinct signal levels



Manchester Phase Encoding

Binary 0 - high to low signal transition

Binary 1 - low to high signal transition

Transition in the center of each bit cell

$0 \rightarrow 1$


$1 \rightarrow 0$

Provides the clock information

Observe

Signal does not return to zero level after each encoded bit

Referred to as *non-return-to-zero NRZ* signal



Manchester Differential Encoding

Transition in the *center* of each bit cell


$0 \rightarrow 1$

$1 \rightarrow 0$

Provides the clock information

Transition at the *start* of each bit cell

Only if next bit to be encoded is a binary 0



Re-derive the Clock

Use a preamble including sync sequence

Done using *phase lock loop- PLL*

Based upon very stable receiver clock

PLL used to keep sample clock locked

Signal transitions of incoming signal

Data encoded to ensure a sufficient number of signal transitions

At each transition sample timing adjusted to ensure

Sampling in center of bit

Scheme will tolerate intervals without transitions

Based upon stability of the clock

Bit / Character / Frame Synchronization

Once individual bits are identified

Potential problems

All devices must run at same speed

Clock skew on long busses


Because of different loading

Have propagation delay along bus

Clock arriving at different times with respect to data

Along bus

Clock noise in system



Potential advantages

Easier to test

Simpler protocol

Easier to stay in sync with data

Interrupts and Polling

Once we have decided upon a the physical configuration

Outside world connection

Several aspects of exchange remain

Will often be interacting with several devices

Must determine

How best to share resources

Cannot afford to continually watch device

To determine if it

Has some information

Is ready to receive (more) information

Several methods by which the is done

Polling

Interrupts

Let's examine each

Polling

With polling scheme

Controller

CPU


Peripheral Processor

In loop

We have the accompanying
algorithm

1. Send device address
2. Receives response
 - The response may consist of a single signal, a word, state information, or status
 - a. If goal to transmit or receive data to device
 - When device ready received
 - Transfer to appropriate routine
 - Execute transfer
 - Return to polling loop
 - b. If goal is to collect status
 - Can do so and continue polling other devices
 - Status may include, self test results, a ready condition after power up, availability for additional transfer, data available.

We express the polling operation in the following state diagram,



In such loop

Control device knows who it is communicating with

Problem with polling

Control device can do nothing else while polling

Advantage however

Process deterministic

Time to complete predictable

Interrupt

Goal of interrupt scheme same as that of polled scheme

Difference

Control device not dedicated to monitoring external devices

I/O system may have from 1 to many interrupt lines

Each line may be connected to 1 or several devices

Single Interrupt Line with Single Device

Let's first examine how single interrupt might work

As name suggests

Interrupt signal interrupts foreground activity

When interrupt occurs several actions possible

Ignore

Respond

Ignore

Will discuss shortly

Respond

Responding to interrupt much like subroutine call

Procedure

Suspend current process

Branch to routine appropriate to interrupting device

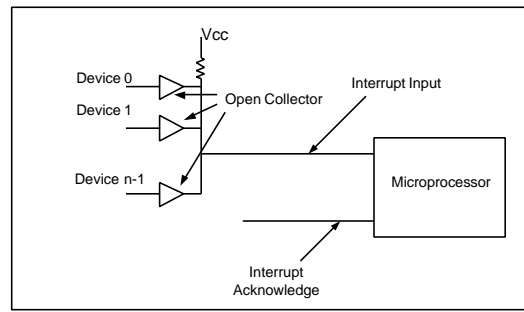
Execute routine

Resume former process

Single Interrupt Line with Multiple Devices

Dealing with multiple devices on single interrupt only slightly more complicated

The interrupt is connected to controller as shown,




Dilemma of course

Identifying which device is requesting service

Observe device priority determined by

Physical proximity to controller

Closest device has highest priority.



Several alternatives

Polling

Vector

Polling

Straight forward

Sequence

Interrupt occurs

Suspend current task

Branch to handler

Poll devices to identify

Branch to appropriate handler routine

Handle

Return

Resume suspended task

Vector

With vector scheme

Device requesting service

Interrupts

When acknowledged

Returns identifier

Device name

Address of service routine

Sequence then proceeds as with polled alternative

Multiple

With several devices on the same line

Must now address question

How to accommodate multiple

Simultaneous interrupts

Sequential interrupts

Simultaneous

Must have priority scheme

Can be

Physical proximity

Assigned

Higher priority device gets serviced first

Sequential

If second interrupt occurs while handling first

Ignore until finished with first

Assign each device priority

If interrupting device has higher priority

Handle in same way

Similar to calling subroutine from subroutine

Multiple Interrupt Lines with Single Device on Each

For each line

Same as single line with single device

With multiple lines

Must assign priority to each

Same as single line with multiple devices

Multiple Interrupt Lines with Multiple Devices on Each

For each line

Same as single line with multiple devices

With multiple lines

Must assign priority to each

Same as single line with multiple devices

